

FEATURE_SELECTION

Filter Methods:

Based on statistical test

Correlation Coefficient: Measures the linear relationship between two variables.

Chi-Square Test: Measures the independence between categorical variables.

Mutual Information: Measures the dependency between variables.

Variance Threshold: Removes features with low variance.

Wrapper Methods:

Forward Selection: Starts with no features, adds one at a time based on model performance.

Backward Elimination: Starts with all features, removes the least significant feature one at a time.

Recursive Feature Elimination (RFE): Recursively removes least important features based on model performance.

Embedded Methods:

Lasso Regression: Adds a penalty equal to the absolute value of the magnitude of coefficients.

Ridge Regression: Adds a penalty equal to the square of the magnitude of coefficients.

Elastic Net: Combines Lasso and Ridge penalties.

Tree-based Methods: Feature importance based on tree-based models like Random Forests and Gradient Boosting.

```
In [14]: import numpy as np
import pandas as pd
from sklearn.feature_selection import VarianceThreshold, mutual_info_classif
from sklearn.linear_model import LogisticRegression, Lasso, Ridge, ElasticNe
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder
from sklearn.impute import KNNImputer
from sklearn.metrics import r2_score
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset
file_path = "D:/Downloads/archive (22)/loan_data_set.csv"
data = pd.read_csv(file_path)

# Drop 'Loan_ID' and separate features and target variable
X = data.drop(['Loan_ID', 'Loan_Status'], axis=1)
y = data['Loan_Status']

# Encode target variable
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

# Encode categorical variables
X_encoded = pd.get_dummies(X)

# Handle missing values
imputer = KNNImputer(n_neighbors=3)
X_imputed = imputer.fit_transform(X_encoded)

# Normalize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)
print(data)
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
..	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
..	
609	2900	0.0	71.0	360.0	
610	4106	0.0	40.0	180.0	
611	8072	240.0	253.0	360.0	
612	7583	0.0	187.0	360.0	
613	4583	0.0	133.0	360.0	

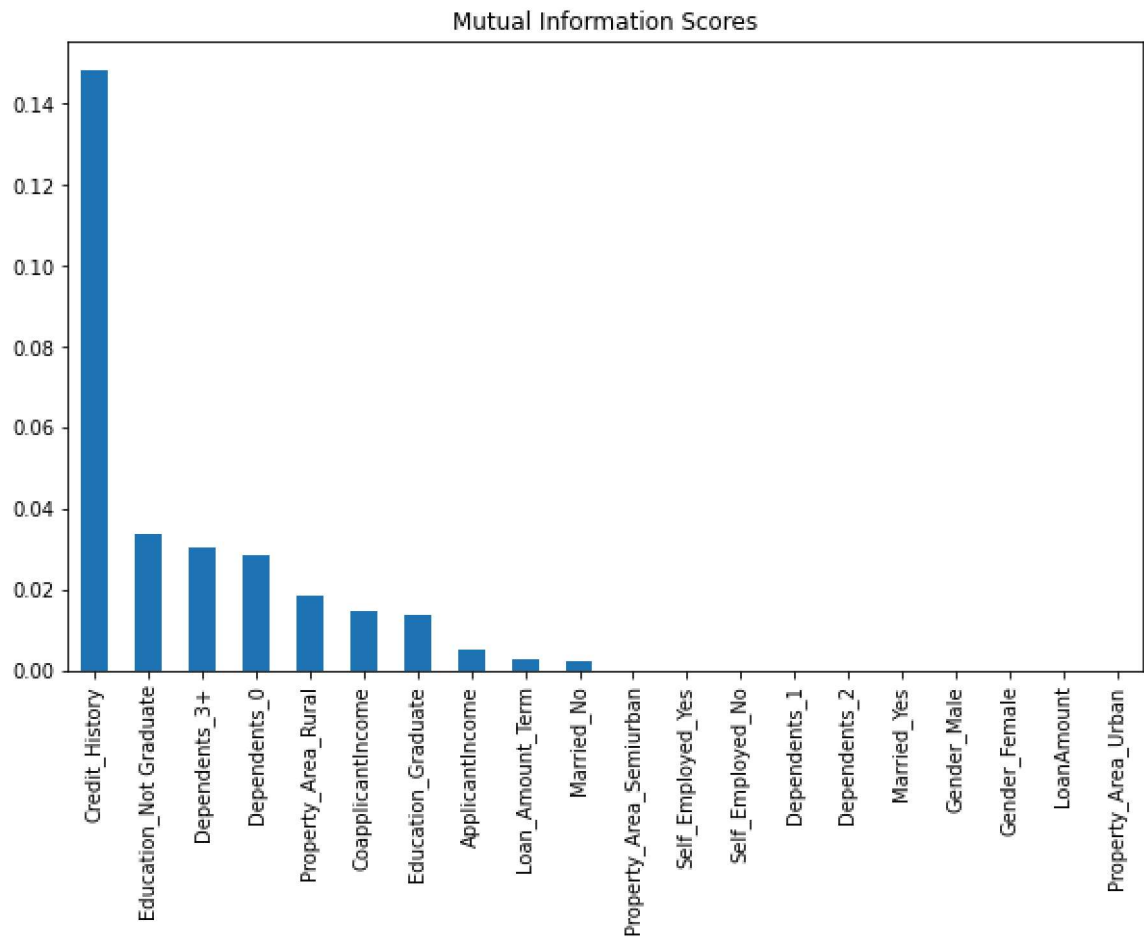
	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y
..
609	1.0	Rural	Y
610	1.0	Rural	Y
611	1.0	Urban	Y
612	1.0	Urban	Y
613	0.0	Semiurban	N

[614 rows x 13 columns]

Filter Methods:

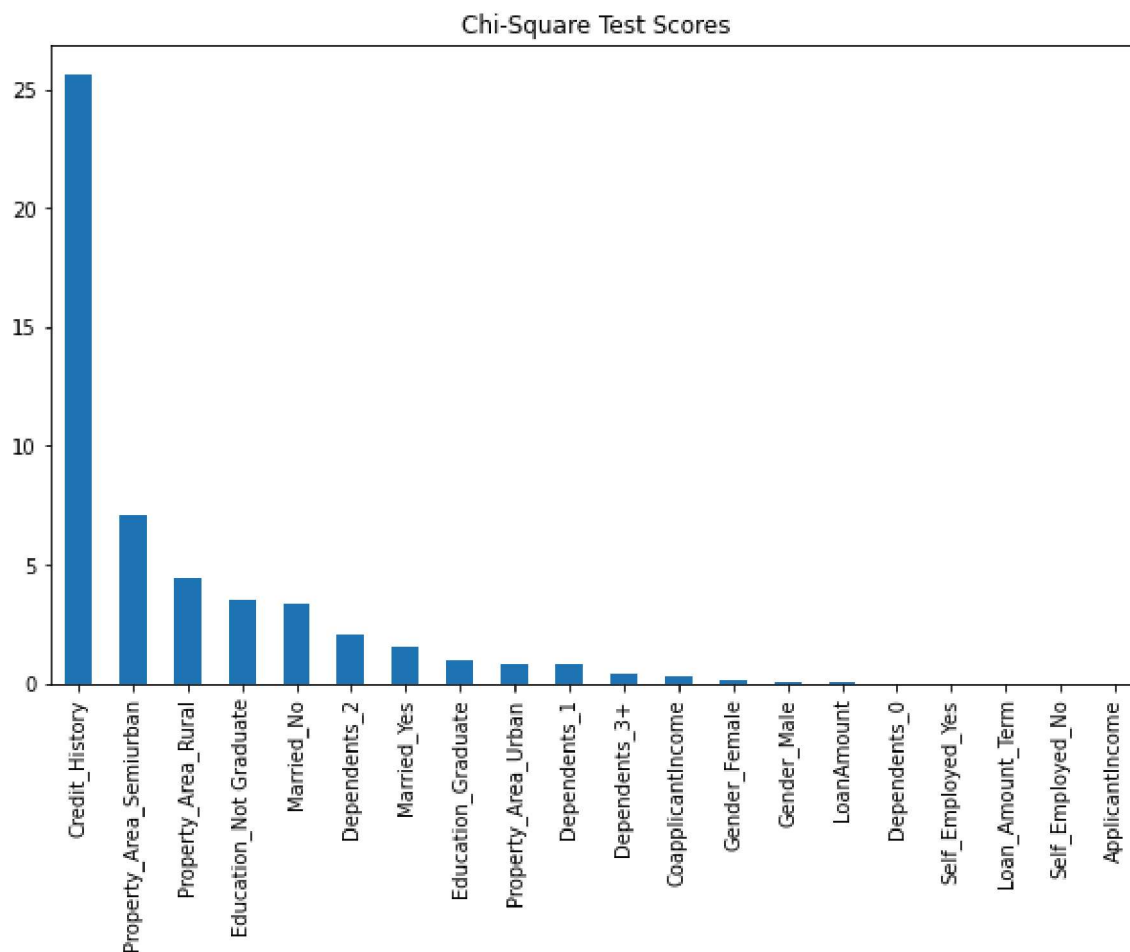
```
In [2]: # Mutual Information
mi = mutual_info_classif(X_scaled, y)
mi_series = pd.Series(mi, index=X_encoded.columns)

plt.figure(figsize=(10, 6))
mi_series.sort_values(ascending=False).plot.bar()
plt.title('Mutual Information Scores')
plt.show()
```



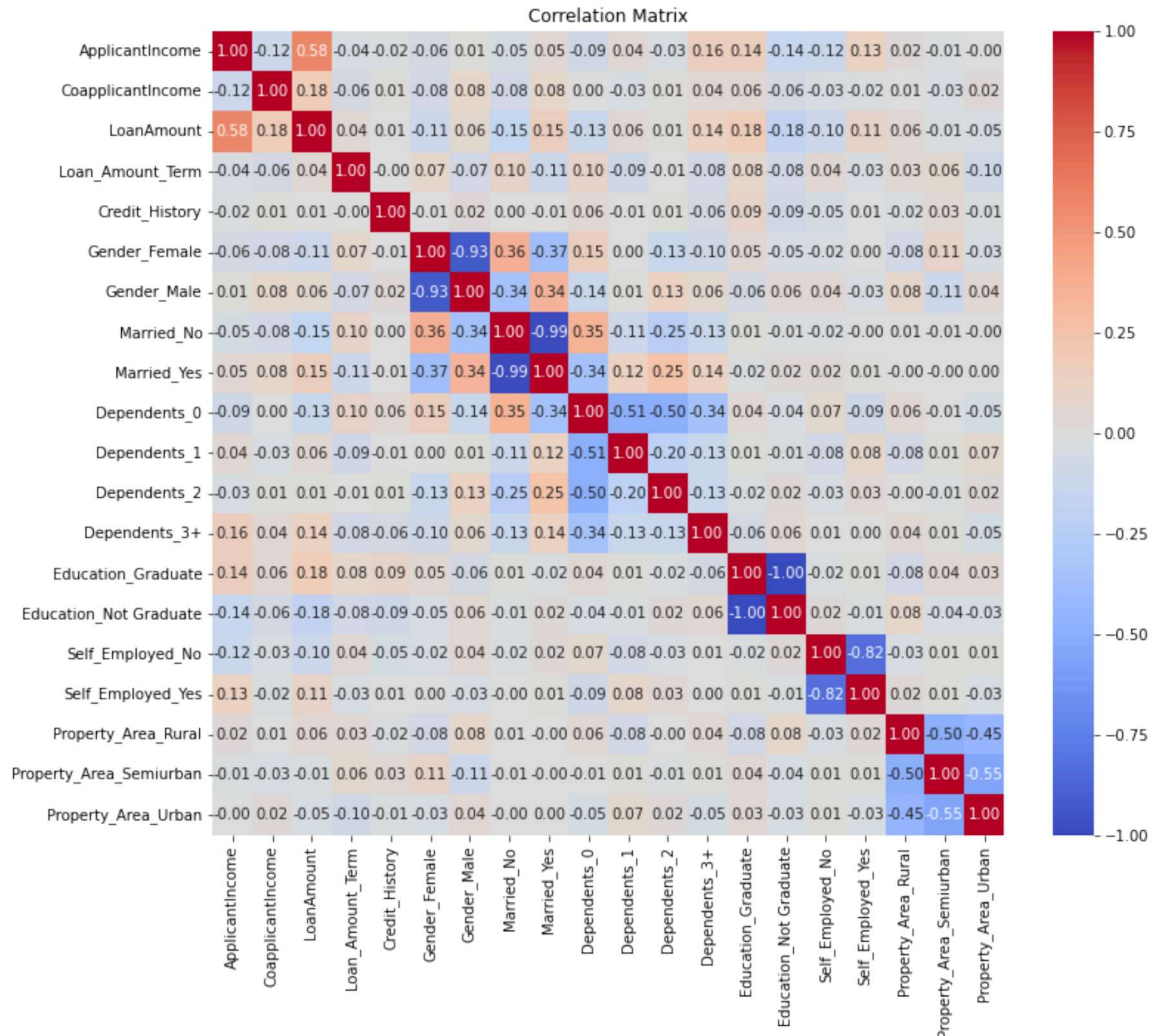
```
In [12]: # Chi-Square Test
min_max_scaler = MinMaxScaler()
X_minmax = min_max_scaler.fit_transform(X_imputed)
chi2_scores, p_values = chi2(X_minmax, y)
chi2_series = pd.Series(chi2_scores, index=X_encoded.columns)

plt.figure(figsize=(10, 6))
chi2_series.sort_values(ascending=False).plot.bar()
plt.title('Chi-Square Test Scores')
plt.show()
```



```
In [15]: # Correlation Coefficient
correlation_matrix = pd.DataFrame(X_imputed, columns=X_encoded.columns).corr

plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```



Wrapper Methods:

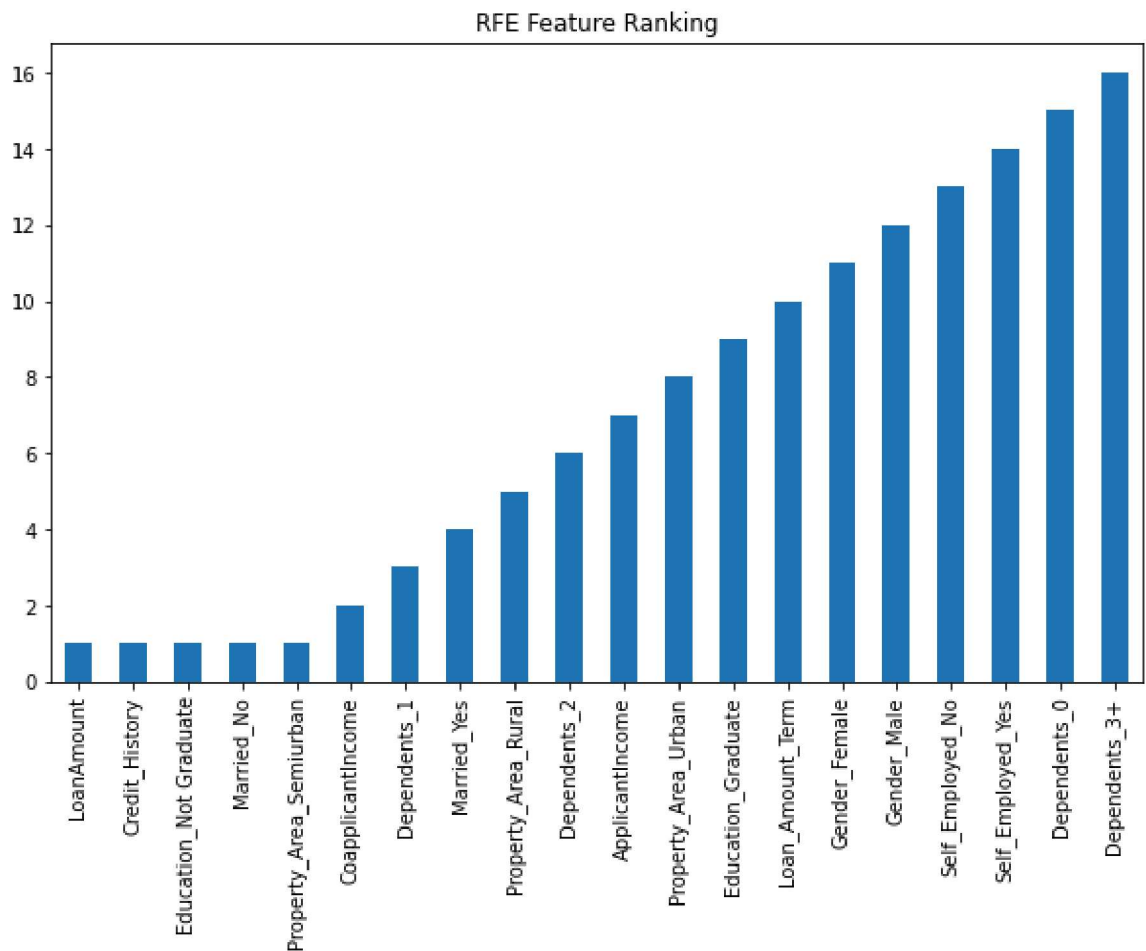
```

In [5]: # Recursive Feature Elimination (RFE)
model = LogisticRegression(max_iter=1000)
rfe = RFE(model, n_features_to_select=5)
X_rfe = rfe.fit_transform(X_scaled, y)
print("RFE Selected Features Shape:", X_rfe.shape)

plt.figure(figsize=(10, 6))
rfe_ranking = pd.Series(rfe.ranking_, index=X_encoded.columns)
rfe_ranking.sort_values().plot.bar()
plt.title('RFE Feature Ranking')
plt.show()

```

RFE Selected Features Shape: (614, 5)



```
In [6]: # Backward Elimination
X_with_const = sm.add_constant(pd.DataFrame(X_scaled, columns=X_encoded.columns))
model = sm.OLS(y, X_with_const).fit()
print("Initial Model Summary:")
print(model.summary())

while True:
    p_values = pd.Series(model.pvalues, index=X_with_const.columns)
    max_p_value = p_values.max()
    if max_p_value > 0.05:
        excluded_feature = p_values.idxmax()
        print(f"Dropping {excluded_feature}")
        X_with_const = X_with_const.drop(columns=[excluded_feature])
        model = sm.OLS(y, X_with_const).fit()
    else:
        break

print("Final Model Summary:")
print(model.summary())

plt.figure(figsize=(10, 6))
p_values.sort_values().plot.bar()
plt.axhline(y=0.05, color='r', linestyle='--')
plt.title('P-Values of Features After Backward Elimination')
plt.show()
```

Initial Model Summary:

OLS Regression Results

```
=====
=====
Dep. Variable:          y      R-squared:
0.319
Model:                OLS      Adj. R-squared:
0.299
Method:             Least Squares      F-statistic:
15.51
Date:                Wed, 05 Jun 2024      Prob (F-statistic):          3.8
9e-39
Time:                10:02:32      Log-Likelihood:          -2
81.12
No. Observations:          614      AIC:
600.2
Df Residuals:              595      BIC:
684.2
Df Model:                  18
```



```

In [7]: # Forward Selection
remaining_features = list(range(X_scaled.shape[1]))
selected_features = []
current_score, best_new_score = 0.0, 0.0

while remaining_features and current_score == best_new_score:
    scores_with_candidates = []
    for candidate in remaining_features:
        features_to_try = selected_features + [candidate]
        model = LinearRegression()
        model.fit(X_scaled[:, features_to_try], y)
        score = r2_score(y, model.predict(X_scaled[:, features_to_try]))
        scores_with_candidates.append((score, candidate))

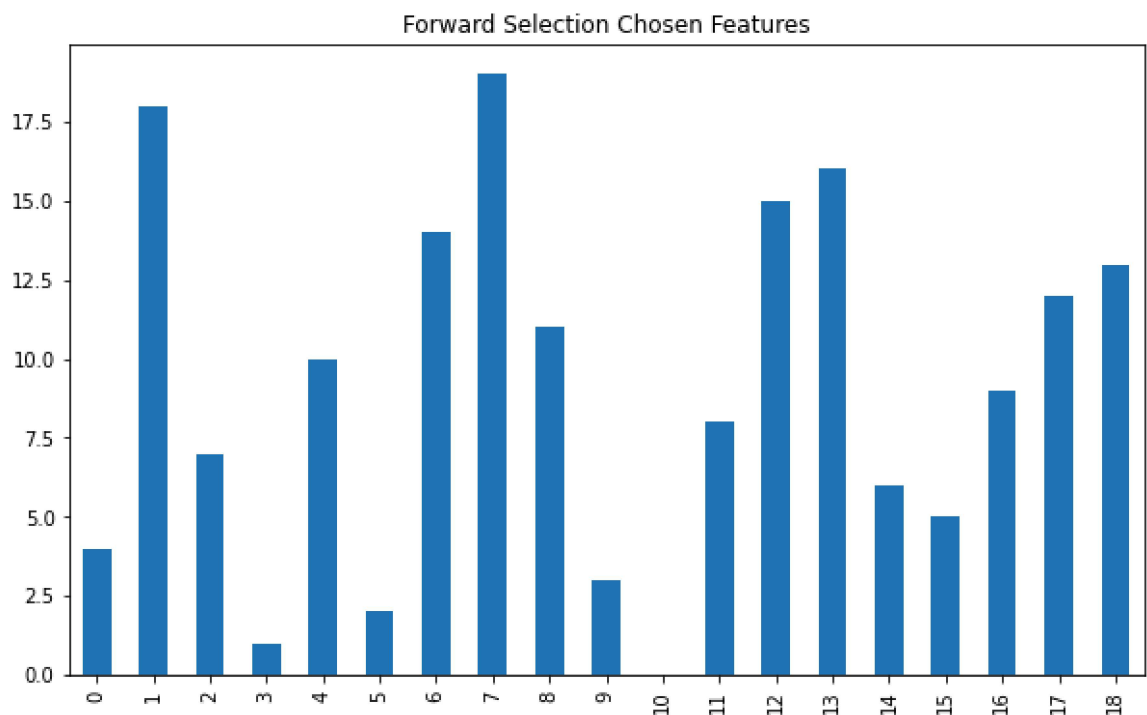
    scores_with_candidates.sort(reverse=True)
    best_new_score, best_candidate = scores_with_candidates[0]
    if current_score < best_new_score:
        remaining_features.remove(best_candidate)
        selected_features.append(best_candidate)
        current_score = best_new_score

print("Selected features (indices):", selected_features)

plt.figure(figsize=(10, 6))
forward_selected_features = pd.Series(selected_features)
forward_selected_features.plot(kind='bar')
plt.title('Forward Selection Chosen Features')
plt.show()

```

Selected features (indices): [4, 18, 7, 1, 10, 2, 14, 19, 11, 3, 0, 8, 15, 16, 6, 5, 9, 12, 13]



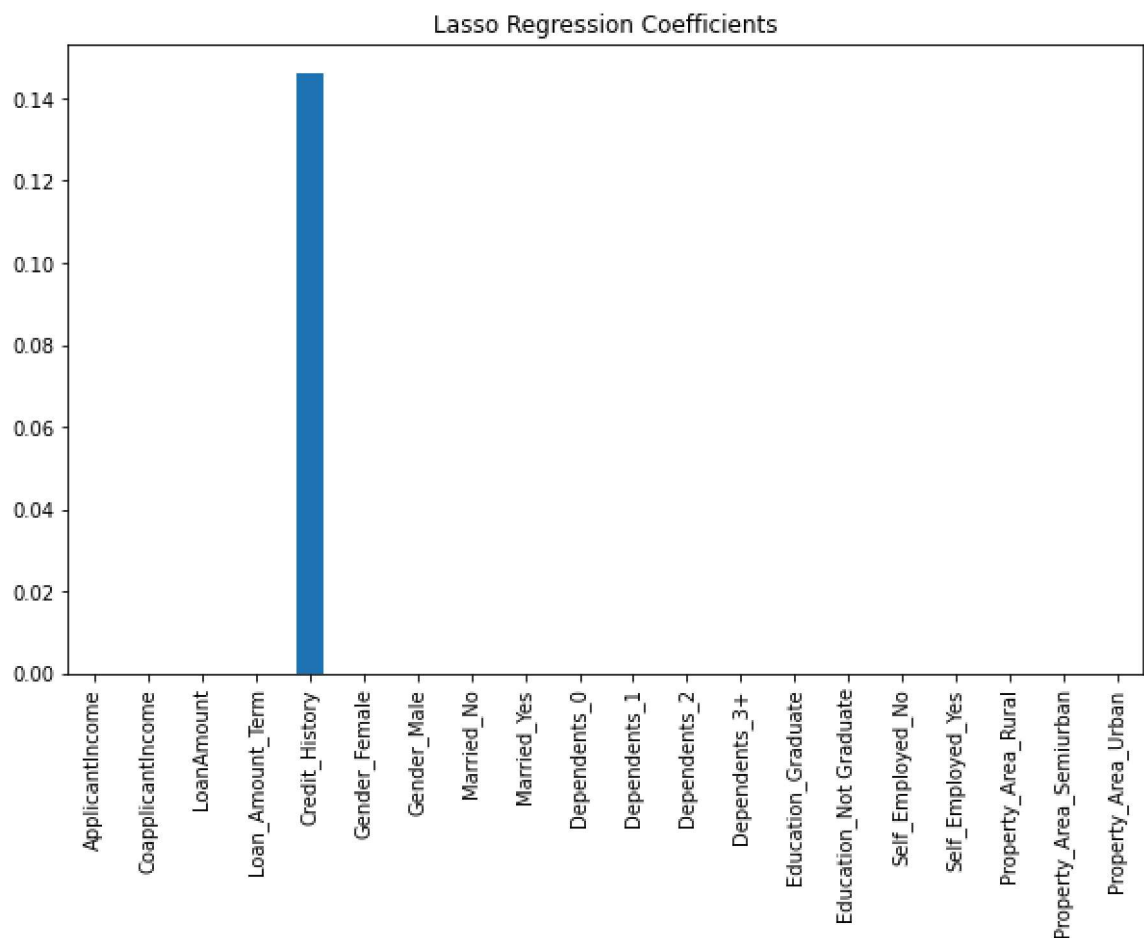
Embedded Methods:

```
In [8]: # Lasso Regression
lasso = Lasso(alpha=0.1)
lasso.fit(X_scaled, y)
print("Lasso Coefficients:")
print(lasso.coef_)

plt.figure(figsize=(10, 6))
lasso_coefficients = pd.Series(lasso.coef_, index=X_encoded.columns)
lasso_coefficients.plot(kind='bar')
plt.title('Lasso Regression Coefficients')
plt.show()
```

Lasso Coefficients:

```
[ 0.         -0.         -0.         -0.         0.14605572 -0.
  0.         -0.         0.         -0.         -0.         0.
 -0.         0.         -0.         0.         -0.         -0.
  0.         -0.         ]
```

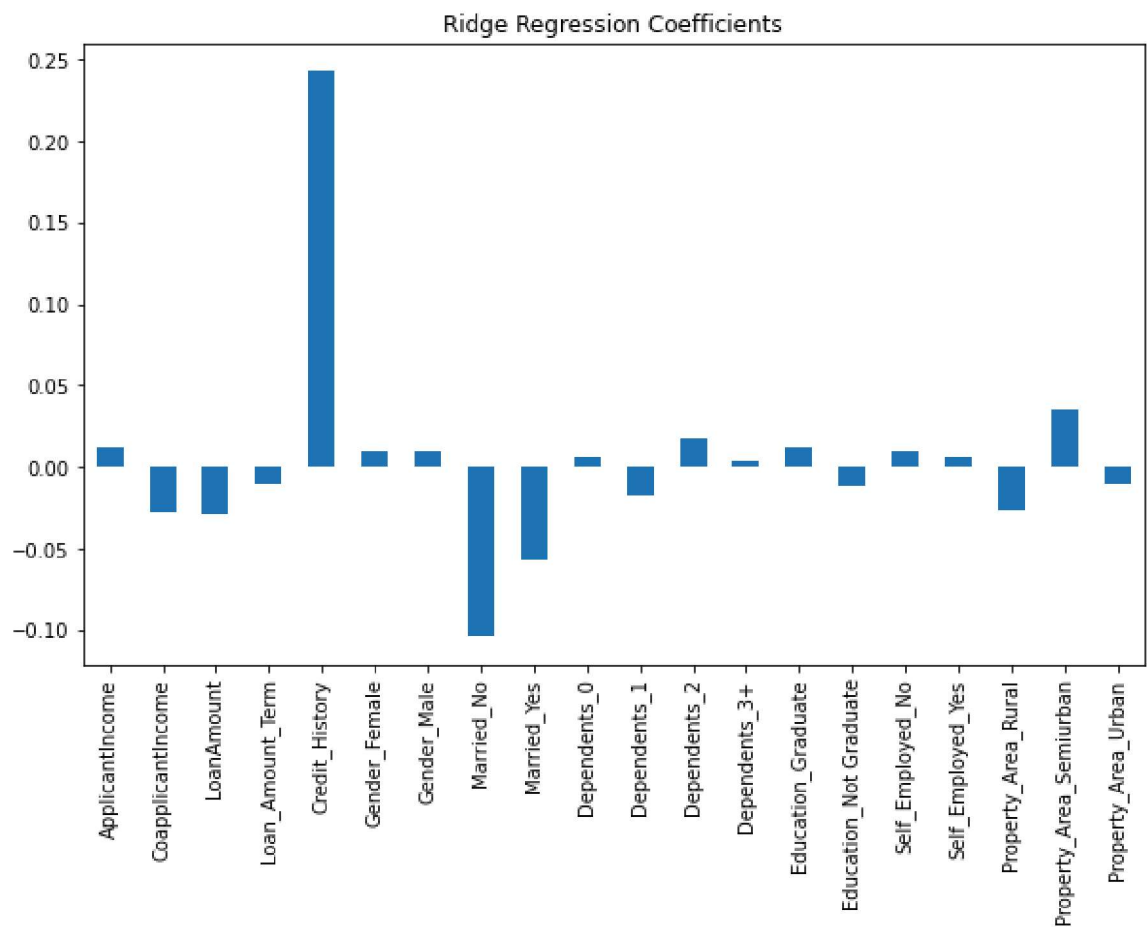


```
In [9]: # Ridge Regression
ridge = Ridge(alpha=0.1)
ridge.fit(X_scaled, y)

plt.figure(figsize=(10, 6))
ridge_coefficients = pd.Series(ridge.coef_, index=X_encoded.columns)
ridge_coefficients.plot(kind='bar')
plt.title('Ridge Regression Coefficients')
plt.show()
```

Ridge Coefficients:

```
[ 0.01198102 -0.02818727 -0.02921305 -0.01001761  0.24270978  0.00996648
 0.00999456 -0.10355527 -0.05721514  0.00577958 -0.01763945  0.01724876
 0.00324774  0.01135585 -0.01135585  0.00914944  0.00659059 -0.0267295
 0.03514268 -0.01044077]
```

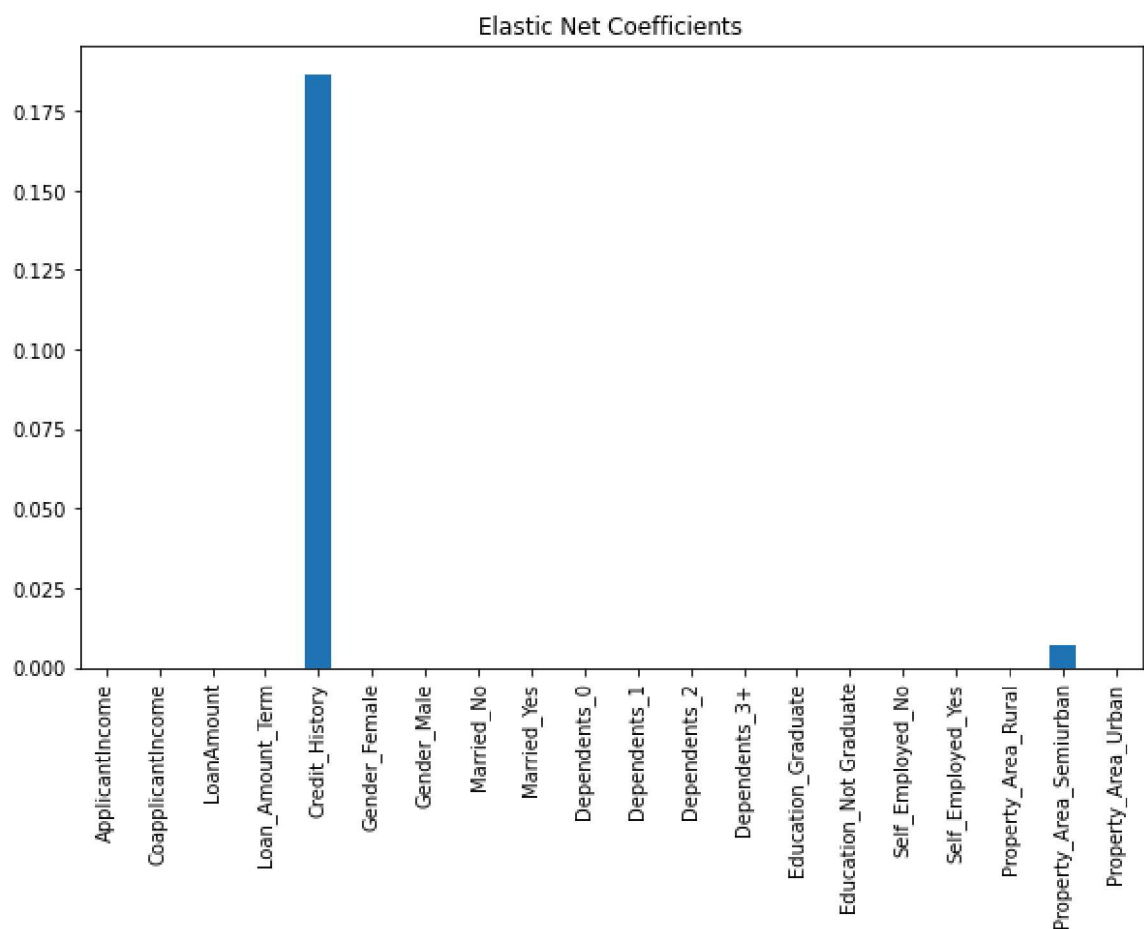


```
In [10]: # Elastic Net
elastic_net = ElasticNet(alpha=0.1, l1_ratio=0.5)
elastic_net.fit(X_scaled, y)

plt.figure(figsize=(10, 6))
elastic_net_coefficients = pd.Series(elastic_net.coef_, index=X_encoded.columns)
elastic_net_coefficients.plot(kind='bar')
plt.title('Elastic Net Coefficients')
plt.show()
```

Elastic Net Coefficients:

```
[ 0.          -0.          -0.          -0.          0.18650781 -0.
  0.          -0.          0.          -0.          -0.          0.
 -0.          0.          -0.          0.          -0.          -0.
  0.00709615 -0.          ]
```



```
In [16]: # Tree-based Methods
model = RandomForestClassifier()
model.fit(X_scaled, y)
importances = model.feature_importances_
importances_series = pd.Series(importances, index=X_encoded.columns)

plt.figure(figsize=(10, 6))
importances_series.sort_values(ascending=False).plot.bar()
plt.title('Random Forest Feature Importances')
plt.show()
```

